

Kernel Debug Zoo

Workshop Edition

Nandor Kracser

<https://github.com/bonifaido/kernel-debug-zoo>

What This Workshop Covers

- Why kernel debugging features matter
- Why you should test against different kernel versions
- Core tools: `lockdep`, `KASAN`, `KMEMLEAK`, `KFENCE`, `SLUB_DEBUG`
- How to trigger them with demo modules
- How Rust could help?

What Is Kernel Debugging?




- Debugging the **Linux kernel** is very different from userspace:
 - No full stack traces by default
 - Panics can crash the system
 - Subtle bugs (e.g., race conditions) hard to reproduce
- Needs special tooling and awareness (out of the box)

 Goal: Add kernel debugging tools to your everyday toolbox

The 2024 CrowdStrike Kernel Panic

- CrowdStrike Falcon Sensor triggered a **kernel panic** on many Linux systems in April 2024
- Cause: faulty eBPF hook in a widely deployed version
- Broke critical infrastructure – DNS, mail, auth, etc.

Lessons:

- Always **test kernel code across versions**
 -  Validate against all upstream kernels
 -  Don't test only the “happy path”
-  Takeaway: kernel-space bugs scale globally.



lockdep

Lock Dependency Validator

- Detects incorrect lock ordering
- Flags possible deadlocks
- Works best when locks are used inconsistently

| WARNING: possible circular locking dependency detected



Usually enabled



Boot param: lockdep



KASAN

Kernel Address Sanitizer

- Detects:
 - Use-after-free
 - Out-of-bounds accesses
- Reports with stack trace



Boot params: `kasan=on kasan.multi_shot=1`



Needs: `CONFIG_KASAN=y`



KFENCE – Kernel Electric Fence

- Lightweight, always-on memory bug detector
- Similar to KASAN in term os features



Tradeoffs:

- Lower accuracy than KASAN
- Near-zero performance overhead
- Safe for **production systems**



Enable: `kfence.sample_interval=100`



Perfect for long-running systems & catching rare bugs



KMEMLEAK

Memory Leak Detector

- Scans memory for unreachable allocations
- Reports leaks via:

`/sys/kernel/debug/kmemleak`



Boot param: `kmemleak=on`

Enable with: `CONFIG_DEBUG_KMEMLEAK`



SLUB Debugging

Heap Hardening


- Detects:
 - Redzone overflows
 - Double frees
 - Use-after-free
- Works with the `kmalloc()` family of functions
- KASAN provides these features



Boot param: `slub_debug=PU`

Page Poisoning

- Catches access to freed pages
- Poisons memory (e.g., with `0xaa`)
- Triggers Oops on access

 Boot param: `page_poison=1`

Needs: `CONFIG_PAGE_POISONING`

Race Condition Demo

- 2 threads update shared var without lock
- Random outputs = race
- `dmesg` might not show it unless it breaks something
- Add manual checks or log divergence

✓ Summary Table

Feature	Catches	Boot Param	Kernel Config
lockdep	Deadlocks	lockdep	CONFIG_LOCKDEP
KASAN	UAF, 00B	kasan=on	CONFIG_KASAN
KMEMLEAK	Leaks	kmemleak=on	CONFIG_DEBUG_KMEMLEAK
SLUB	Heap overflows	slub_debug=PU	CONFIG_SLUB_DEBUG_ON
Page Poison	Freed-page access	page_poison=1	CONFIG_PAGE_POISONING

Tips for Demos

- Trigger bugs safely: prefer VMs
- Run `dmesg -w` while testing
- Use `pr_info()` or `trace_printk()`
- Clean up modules: `rmmod`

Enabling These Features

Option 1: Use a Prebuilt Debug Kernel

- Fedora/CentOS:

```
sudo dnf install kernel-debug sudo grubby --set-default /boot/vmlinuz-<debug-version>
```
- Comes with most debug configs: `KASAN` , `lockdep` , `SLUB_DEBUG` , etc.

Option 2: Build a Custom Kernel

- Set config options in `make menuconfig` :
- `CONFIG_KASAN=y`
- `CONFIG_DEBUG_KMEMLEAK=y`
- `CONFIG_LOCKDEP_SUPPORT=y`
- `CONFIG_SLUB_DEBUG_ON=y`
- `CONFIG_PAGE_POISONING=y`
- Add boot params in GRUB:
`kasan=on slub_debug=PU kmemleak=on page_poison=1`

Can Rust Help?

Yes – a lot.

- Rust prevents **entire classes** of memory bugs:
- Use-after-free
- Buffer overflows
- Null pointer deref
- Enforced at compile-time

Rust + Kernel = Safer Drivers

- Rust kernel modules have:
- **Ownership model** for memory safety
- Safe concurrency via `Send` / `Sync`
- No `unsafe` code unless explicitly declared

 Still early-stage, but promising for:

- Drivers
- Filesystems
- Net modules

Rust vs C (for Kernel Safety)

Bug Type	C Kernel	Rust Kernel
Use-after-free	✓	⊘ (safe by default)
Buffer overflow	✓	⊘
Null deref	✓	⊘
Data races	✓	⊘ (unless <code>unsafe</code>)

 **Questions?**

Let's go break stuff. 